



Object Detection

In this assignment, you will develop an object detector based on gradient features and sliding window classification. A set of test images and *hogvis.py* are provided in the Canvas assignment directory

Name: Kei Asakawa

SID: 22659423

```
In [1]: 1 import numpy as np
        2 import matplotlib.pyplot as plt
```

1. Image Gradients [20 pts]

Write a function that takes a grayscale image as input and returns two arrays the same size as the image, the first of which contains the magnitude of the image gradient at each pixel and the second containing the orientation.

Your function should filter the image with the simple x- and y-derivative filters described in class. Once you have the derivatives you can compute the orientation and magnitude of the gradient vector at each pixel. You should use ***scipy.ndimage.correlate*** with the 'nearest' option in order to nicely handle the image boundaries.

Include a visualization of the output of your gradient calculate for a small test image. For displaying the orientation result, please uses a cyclic colormap such as "hsv" or "twilight". (see <https://matplotlib.org/tutorials/colors/colormaps.html> (<https://matplotlib.org/tutorials/colors/colormaps.html>))

NOTE: To be consistent with the provided code that follows, the gradient orientation values you return should range in $(-\pi/2, +\pi/2)$ where a horizontal edge (vertical gradient) is $-\pi/2$ and the angle increases as the edge rotates clockwise in the image.


In [2]:


```
1  #we will only use:  scipy.ndimage.correlate
2  from scipy import ndimage
3
4  def mygradient(image):
5      """
6      This function takes a grayscale image and returns two arrays of the
7      same size, one containing the magnitude of the gradient, the second
8      containing the orientation of the gradient.
9
10
11     Parameters
12     -----
13     image : 2D float array of shape HxW
14             An array containing pixel brightness values
15
16     Returns
17     -----
18     mag : 2D float array of shape HxW
19           gradient magnitudes
20
21     ori : 2Dfloat array of shape HxW
22           gradient orientations in radians
23     """
24     h = image.shape[0]
25     w = image.shape[1]
26     mag = np.zeros((h,w))
27     ori = np.zeros((h,w))
28     w_x = [[0,1]]
29     w_y = [[0],
30            [1]]
31     dy = np.zeros((h,w))
32     dx = np.zeros((h,w))
33     for x in range(h):
34         for y in range(w):
35             if x != h - 1:
36                 dy[x,y] = image[x+1, y] - image[x, y]
37             else:
38                 dy[x,y] = ndimage.correlate(dy, w_y, mode='nearest')[x,y]
39             if y != w - 1:
40                 dx[x,y] = image[x, y+1] - image[x, y]
41             else:
42                 dx[x,y] = ndimage.correlate(dx, w_x, mode='nearest')[x,y]
```

```
43         ori[x,y] = np.arctan(dy[x,y] / (dx[x,y] + 1e-4))
44         mag[x,y] = np.sqrt(np.square(dx[x,y]) + np.square(dy[x,y]))
45     return (mag,ori)
```

In [3]:



```
1 #
2 # Demonstrate your mygradient function here by loading in a grayscale
3 # image, calling mygradient, and visualizing the resulting magnitude
4 # and orientation images. For visualizing orientation image, I suggest
5 # using the hsv or twilight colormap.
6 #
7
8 # here is one simple test image which has gradients pointed in all
9 # directions so you can see if your orientation estimates are reasonable
10 [yy,xx] = np.mgrid[-100:100,-100:100]
11 testimage = np.minimum(np.maximum(np.array(xx*xx+yy*yy,dtype=float),400),8100)
12
13
14 # you should also load in or synthesize another image to test with besides
15 # the one above.
16
17 (mag,ori) = mygradient(testimage)
18
19 #visualize results mag,ori as images
20 plt.imshow(mag, cmap='twilight')
21 plt.show()
22 plt.imshow(ori, cmap='twilight')
23 plt.show()
24
```





Input Image

Gradient Vector



Cell Histogram



HOG Features

2. Histograms of Gradient Orientations [25 pts]

Write a function that computes gradient orientation histograms over each 8x8 block of pixels in an image. Your function should bin the orientation into 9 equal sized bins between $-\pi/2$ and $\pi/2$. The input of your function will be an image of size HxW. The output should be a three-dimensional array **ohist** whose size is $(H/8) \times (W/8) \times 9$ where **ohist[i,j,k]** contains the count of how many edges of orientation k fell in block (i,j). If the input image dimensions are not a multiple of 8, you should use **np.pad** with the **mode=edge** option to pad the width and height up to the nearest integer multiple of 8.

To determine if a pixel is an edge, we need to choose some threshold. I suggest using a threshold that is 10% of the maximum gradient magnitude in the image. Since each 8x8 block will contain a different number of edges, you should normalize the resulting histogram for each block to sum to 1 (i.e., **np.sum(ohist,axis=2)** should be 1 at every location).

I would suggest your function loops over the orientation bins. For each orientation bin you'll need to identify those pixels in the image whose gradient magnitude is above the threshold and whose orientation falls in the given bin. You can do this easily in numpy using logical operations in order to generate an array the same size as the image that contains Trues at the locations of every edge pixel that falls in the given orientation bin and is above threshold. To collect up pixels in each 8x8 spatial block you can use the function **ski.util.view_as_windows(...,(8,8),step=8)** and **np.count_nonzeros** to count the number of edges in each block.

Test your code by creating a simple test image (e.g. a white disk on a black background), computing the descriptor and using the provided function **hogvis** to visualize it.

Note: in the discussion above I have assumed 8x8 block size and 9 orientations. In your code you should use the parameters **bsize** and **noorient** in place of these constants.

In [4]:

```
1  #we will only use: ski.util.view_as_windows for computing hog descriptor
2  import skimage as ski
3
4  def hog(image, bsize=8, norient=9):
5
6      """
7      This function takes a grayscale image and returns a 3D array
8      containing the histogram of gradient orientations descriptor (HOG)
9      We follow the convention that the histogram covers gradients starting
10     with the first bin at -pi/2 and the last bin ending at pi/2.
11
12     Parameters
13     -----
14     image : 2D float array of shape HxW
15            An array containing pixel brightness values
16
17     bsize : int
18            The size of the spatial bins in pixels, defaults to 8
19
20     norient : int
21            The number of orientation histogram bins, defaults to 9
22
23     Returns
24     -----
25     ohist : 3D float array of shape (H/bsize,W/bsize,norient)
26            edge orientation histogram
27
28     """
29
30     # determine the size of the HOG descriptor
31     (h,w) = image.shape
32     h2 = int(np.ceil(h/float(bsize)))
33     w2 = int(np.ceil(w/float(bsize)))
34     ohist = np.zeros((h2,w2,norient))
35
36     # pad the input image on right and bottom as needed so that it
37     # is a multiple of bsize
38     if (w2*bsize - w) % 2 == 0:
39         pw = ((w2*bsize-w) // 2, (w2*bsize-w) // 2)
40     else:
41         pw = ((w2*bsize-w) // 2, (w2*bsize-w) // 2 + 1) #amounts to pad on left and right side
42     if (h2*bsize - h) % 2 == 0:
```



```

43     ph = ((h2*bsize-h) // 2, (h2*bsize-h) // 2)
44 else:
45     ph = ((h2*bsize-h) // 2, (h2*bsize-h) // 2 + 1) #amounts to pad on bottom and top side
46 image = np.pad(image,(ph, pw), mode='edge')
47
48 # make sure we did the padding correctly
49 assert(image.shape==(h2*bsize,w2*bsize))
50
51 # compute image gradients
52 (mag,ori) = mygradient(image)
53
54 # choose a threshold which is 10% of the maximum gradient magnitude in the image
55 thresh = mag[np.unravel_index(np.argmax(mag, axis=None), mag.shape)] * .1
56
57
58 # separate out pixels into orientation channels, dividing the range of orientations
59 # [-pi/2,pi/2] into norient equal sized bins and count how many fall in each block
60 binEdges = np.linspace(-np.pi/2, np.pi/2, norient+1);
61 # as a sanity check, make sure every pixel gets assigned to at most 1 bin.
62 bincount = np.zeros((h2*bsize,w2*bsize))
63 for i in range(norient):
64     #create a binary image containing 1s for pixels at the ith
65     #orientation where the magnitude is above the threshold.
66     if i != norient - 1:
67         B = np.logical_and(np.logical_and(binEdges[i + 1] > ori, ori >= binEdges[i]), mag >= thresh)
68     else:
69         B = np.logical_and(ori >= binEdges[i], mag >= thresh)
70     #sanity check: record which pixels have been selected at this orientation
71     bincount = bincount + B
72
73     #pull out non-overlapping bsize x bsize blocks
74     chblock = ski.util.view_as_windows(B,(bsize,bsize),step=bsize)
75
76     #sum up the count for each block and store the results
77     sums = np.zeros((chblock.shape[0], chblock.shape[1]))
78     for j in range(len(chblock)):
79         for k in range(len(chblock[j])):
80             sums[j,k] = np.count_nonzero(chblock[j,k])
81     ohist[:, :, i] = sums
82 #each pixel should have only selected at most once
83 assert(np.all(bincount<=1))
84
85 # Lastly, normalize the histogram so that the sum along the orientation dimension is 1


```

```
86     # note: don't divide by 0! If there are no edges in a block (i.e. the sum of counts  
87 # is 0) then your code should leave all the values as zero.  
88     s = np.sum(ohist, axis = 2)  
89     s = np.where(s == 0, 1, s)  
90     for i in range(norient):  
91         ohist[:, :, i] /= s  
92  
93  
94     assert(ohist.shape==(h2,w2,norient))  
95  
96     return ohist
```

In [5]:

```
1  #provided function for visualizing hog descriptors
2  from hogvis import hogvis
3
4  #
5  # generate a simple test image... a 80x80 image
6  # with a circle of radius 30 in the center
7  #
8  [yy,xx] = np.mgrid[-40:41,-40:41]
9  im = np.array((xx*xx+yy*yy<=30*30),dtype=float)
10 #
11 # display the image and the output of hogvis
12 #
13 hogim = hogvis(hog(im))
14
15 plt.subplot(1,2,1)
16 plt.imshow(im)
17 plt.subplot(1,2,2)
18 plt.imshow(hogim)
19 plt.show()
20
21 # two other synthetic test images to experiment with
22 [yy,xx] = np.mgrid[-40:41,-40:41]
23 im = np.array((xx*xx+yy*yy<=30*30),dtype=float)
24 im[np.abs(xx+yy)<=3] = 0
25
26 hogim = hogvis(hog(im))
27
28 plt.subplot(1,2,1)
29 plt.imshow(im)
30 plt.subplot(1,2,2)
31 plt.imshow(hogim)
32 plt.show()
33
34 [yy,xx] = np.mgrid[-100:101,-100:101]
35 im = np.array(np.sin((xx*xx+yy*yy)/800),dtype=float)
36
37 hogim = hogvis(hog(im))
38
39 plt.subplot(1,2,1)
40 plt.imshow(im)
41 plt.subplot(1,2,2)
42 plt.imshow(hogim)
```

```
43 plt.show()  
44
```



3. Detection [25 pts]

Write a function that takes a template and an image and returns the top detections found in the image. Your function should follow the definition given below.

In your function you should first compute the histogram-of-gradient-orientation feature map for the image, then correlate the template with the feature map. Since the feature map and template are both three dimensional, you will want to filter each orientation separately and then sum up the results to get the final response. If the image of size $H \times W$ then this final response map will be of size $(H/8) \times (W/8)$.

When constructing the list of top detections, your code should implement non-maxima suppression so that it doesn't return overlapping detections. You can do this by sorting the responses in descending order of their score. Every time you add a detection to the list to return, check to make sure that the location of this detection is not too close to any of the detections already in the output list. You can estimate the overlap by computing the distance between a pair of detections and checking that the distance is greater than say 70% of the width of the template.

Your code should return the locations of the detections in terms of the original image pixel coordinates (so if your detector had a high response at block $[i, j]$ in the response map, then you should return $(8i, 8j)$ as the pixel coordinates).

I have provided a function for visualizing the resulting detections which you can use to test your detect function. Please include some visualization of a simple test case.

In [6]:

```
1  #we will only use: scipy.ndimage.correlate
2  from scipy import ndimage
3
4  def detect(image,template,ndetect=5,bsize=8,norient=9):
5
6      """
7      This function takes a grayscale image and a HOG template and
8      returns a list of detections where each detection consists
9      of a tuple containing the coordinates and score (x,y,score)
10
11      Parameters
12      -----
13      image : 2D float array of shape HxW
14             An array containing pixel brightness values
15
16      template : a 3D float array
17                 The HOG template we wish to match to the image
18
19      ndetect : int
20                 Maximum number of detections to return
21
22      bsize : int
23                 The size of the spatial bins in pixels, defaults to 8
24
25      norient : int
26                 The number of orientation histogram bins, defaults to 9
27
28      Returns
29      -----
30      detections : a list of tuples of length ndetect
31                   Each detection is a tuple (x,y,score)
32
33      """
34
35      # norient for the template should match the norient parameter passed in
36      assert(template.shape[2]==norient)
37
38      fmap = hog(image,bsize=bsize,norient=norient)
39
40
41      #cross-correlate the template with the feature map to get the total response
42      resp = np.zeros((fmap.shape[0],fmap.shape[1]))
```

```

43 for i in range(norient):
44     resp = resp + ndimage.correlate(fmap[:, :, i], template[:, :, i])
45
46     #sort the values in resp in descending order.
47     # val[i] should be ith largest score in resp
48     # ind[i] should be the index at which it occurred so that val[i]==resp[ind[i]]
49     #
50     f = np.ndarray.flatten(resp)
51     val = np.flip(np.sort(f)) #sorted response values
52     ind = np.flip(np.unravel_index(np.argsort(resp, axis=None), resp.shape)) #corresponding indices
53
54     #work down the list of responses from high to low, to generate a
55     # list of ndetect top scoring matches which do not overlap
56     detcount = 0
57     i = 0
58     detections = []
59     while ((detcount < ndetect) and (i < len(val))):
60         # convert 1d index into 2d index
61         yb = ind[1][i]
62         xb = ind[0][i]
63
64         assert(val[i]==resp[yb,xb]) #make sure we did indexing correctly
65
66         #covert block index to pixel coordinates based on bsize
67         xp = bsize*xb
68         yp = bsize*yb
69
70         #check if this detection overlaps any detections that we've already added
71         #to the list. compare the x,y coordinates of this detection to the x,y
72         #coordinates of the detections already in the list and see if any overlap
73         #by checking if the distance between them is less than 70% of the template
74         # width/height
75         overlap = False
76         for j in range(len(detections)):
77             if np.sqrt(np.square(xp-detections[j][0]) + np.square(yp-detections[j][1])) < template.shape[1]
78                 overlap = True
79
80
81         #if the detection doesn't overlap then add it to the list
82         if not overlap:
83             detcount = detcount + 1
84             detections.append((xp,yp,val[i]))
85

```

```
86         i=i+1
87
88     if (len(detections) < ndetect):
89         print('WARNING: unable to find ',ndetect,' non-overlapping detections')
90
91     return detections
```



In [7]:

```
1 import matplotlib.patches as patches
2
3 def plot_detections(image,detections,tsize_pix):
4     """
5     This is a utility function for visualization that takes an image and
6     a list of detections and plots the detections overlayed on the image
7     as boxes.
8
9     Color of the bounding box is based on the order of the detection in
10    the list, fading from green to red.
11
12    Parameters
13    -----
14    image : 2D float array of shape HxW
15           An array containing pixel brightness values
16
17    detections : a list of tuples of length ndetect
18                Detections are tuples (x,y,score)
19
20    tsize_pix : (int,int)
21               The height and width of the box in pixels
22
23    Returns
24    -----
25    None
26
27    """
28    ndetections = len(detections)
29
30    plt.imshow(image,cmap=plt.cm.gray)
31    ax = plt.gca()
32    w = tsize_pix[1]
33    h = tsize_pix[0]
34    red = np.array([1,0,0])
35    green = np.array([0,1,0])
36    ct = 0
37    for (x,y,score) in detections:
38        xc = x-(w//2)
39        yc = y-(h//2)
40        col = (ct/ndetections)*red + (1-(ct/ndetections))*green
41        rect = patches.Rectangle((xc,yc),w,h,linewidth=3,edgecolor=col,facecolor='none')
42        ax.add_patch(rect)
```

```
43         ct = ct + 1
44
45     plt.show()
```

In [8]:

```
1 #
2 # sketch of some simple test code, modify as needed
3 #
4
5 #create a synthetic image with some overlapping circles
6 [yy,xx] = np.mgrid[-40:40,-80:80]
7 im1 = np.array((xx*xx+yy*yy<=30*30),dtype=float)
8 [yy,xx] = np.mgrid[-40:40,-40:120]
9 im2 = np.array((xx*xx+yy*yy<=25*25),dtype=float)
10 [yy,xx] = np.mgrid[-40:40,-120:40]
11 im3 = np.array((xx*xx+yy*yy<=25*25),dtype=float)
12 im = (1/3)*(im1+im2+im3)
13
14
15 #compute feature map with default parameters
16 fmap = hog(im)
17
18 #extract a 3x3 template
19 template = fmap[2:5,2:5,:]
20
21 #run the detect code
22 detections = detect(im,template,ndetect=5)
23
24 #visualize results.
25 plot_detections(im,detections,(3*8,3*8))
26
27 # visually confirm that:
28 # 1. top detection should be the same as the location where we selected the template
29 # 2. multiple detections do not overlap too much
```



4. Learning Templates [15 pts]

The final step is to implement a function to learn a template from positive and negative examples. Your code should take a collection of cropped positive and negative examples of the object you are interested in detecting, extract the features for each, and generate a template by taking the average positive template minus the average negative template.

In [16]:

```
1 import skimage.transform
2 def learn_template(posfiles,negfiles,tsize=np.array([16,16]),bsize=8,norient=9):
3     """
4     This function takes a list of positive images that contain cropped
5     examples of an object + negative files containing cropped background
6     and a template size. It produces a HOG template and generates visualization
7     of the examples and template
8
9     Parameters
10    -----
11    posfiles : list of str
12               Image files containing cropped positive examples
13
14    negfiles : list of str
15               Image files containing cropped negative examples
16
17    tsize : (int,int)
18            The height and width of the template in blocks
19
20    Returns
21    -----
22    template : float array of size tsize x norient
23               The learned HOG template
24
25    """
26
27    #compute the template size in pixels
28    #corresponding to the specified template size (given in blocks)
29    tsize_pix=bsize*tsize
30
31    #figure to show positive training examples
32    fig1 = plt.figure()
33    pltct = 1
34
35    #accumulate average positive and negative templates
36    pos_t = np.zeros((tsize[0],tsize[1],norient),dtype=float)
37    for file in posfiles:
38        #load in a cropped positive example
39        img = plt.imread(file)
40
41        #convert to grayscale and resize to fixed dimension tsize_pix
42        #using skimage.transform.resize if needed.
```

```

43     img = np.mean(img, axis=2)
44     img_scaled = skimage.transform.resize(img, (tsize_pix))
45
46     #display the example. if you want to train with a large # of examples,
47     #you may want to modify this, e.g. to show only the first 5.
48     ax = fig1.add_subplot(len(posfiles),1,pltct)
49     ax.imshow(img_scaled,cmap=plt.cm.gray)
50     pltct = pltct + 1
51
52     #extract feature
53     fmap = hog(img_scaled, bsize, norient)
54     hogim = hogvis(fmap)
55
56     plt.imshow(hogim)
57     plt.show()
58
59     #compute running average
60     pos_t = np.add(pos_t, fmap)
61
62     pos_t = (1/len(posfiles))*pos_t
63     plt.show()
64
65     # repeat same process for negative examples
66     fig2 = plt.figure()
67     pltct = 1
68     neg_t = np.zeros((tsize[0],tsize[1],norient),dtype=float)
69     for file in negfiles:
70         img = plt.imread(file)
71         img = np.mean(img, axis=2)
72         img_scaled = skimage.transform.resize(img, (tsize_pix))
73         ax = fig2.add_subplot(len(negfiles),1,pltct)
74         ax.imshow(img_scaled,cmap=plt.cm.gray)
75         pltct = pltct + 1
76         fmap = hog(img_scaled, bsize, norient)
77         neg_t = np.add(neg_t, fmap)
78
79
80     neg_t = (1/len(negfiles))*neg_t
81     plt.show()
82
83     # add code here to visualize the positive and negative parts of the template
84     # using hogvis. you should separately visualize pos_t and neg_t rather than
85     # the final tempalte.

```

```

86     hogim = hogvis(pos_t)
87
88     plt.imshow(hogim)
89     plt.show()
90
91     hogim = hogvis(neg_t)
92
93     plt.imshow(hogim)
94     plt.show()
95
96     # now construct our template as the average positive minus average negative
97     template = pos_t - neg_t
98
99
100    return template
101

```

In []:

1

5. Experiments [15 pts]

Test your detection by training a template and running it on a test image.

In your experiments and writeup below you should include: (a) a visualization of the positive and negative patches you use to train the template and corresponding hog feature, (b) the detection results on the test image. You should show (a) and (b) for **two different object categories**, the provided face test images and another category of your choosing (e.g. feel free to experiment with detecting cat faces, hands, cups, chairs or some other type of object). Additionally, please include results of testing your detector where there are at least 3 objects to detect (this could be either 3 test images which each have one or more objects, or a single image with many (more than 3) objects). Your test image(s) should be distinct from your training examples. Finally, write a brief (1 paragraph) discussion of where the detector works well and when it fails. Describe some ways you might be able to make it better.



NOTE 1: You will need to create the cropped test examples to pass to your **learn_template**. You can do this by cropping out the examples by hand (e.g. using an image editing tool). You should attempt to crop them out in the most consistent way possible, making sure that each example is centered with the same size and aspect ratio. Negative examples can be image patches that don't contain the object of interest. You should crop out negative examples with roughly the same resolution as the positive examples.



NOTE 2: For the best result, you will want to test on images where the object is the same size as your template. I recommend using the default ***bsize*** and ***noorient*** parameters for all your experiments. You will likely want to modify the template size as needed



Experiment 1: Face detection



In [17]:


```
1 # assume template is 16x16 blocks, you may want to adjust this
2 # for objects of different size or aspect ratio.
3 # compute image a template size
4 bsize=8
5 tsize=np.array([15,15]) #height and width in blocks
6 tsize_pix = bsize*tsize #height and width in pixels
7 posfiles = ('pos1.jpg','pos2.jpg','pos3.jpg','pos4.jpg','pos5.jpg')
8 negfiles = ('neg1.jpg','neg2.jpg','neg3.jpg','neg4.jpg','neg5.jpg')
9
10 # call learn_template to learn and visualize the template and training data
11 t = learn_template(posfiles,negfiles,tsize=tsize)
12
13 # call detect on one or more test images, visualizing the result with the plot_detections function
14 im = plt.imread("images/faces/faces4.jpg")
15 im = np.mean(im, axis=2)
16
17 detections = detect(im, t, 7)
18 plot_detections(im,detections,tsize_pix)
19
```

















Experiment 2: Cat face detection

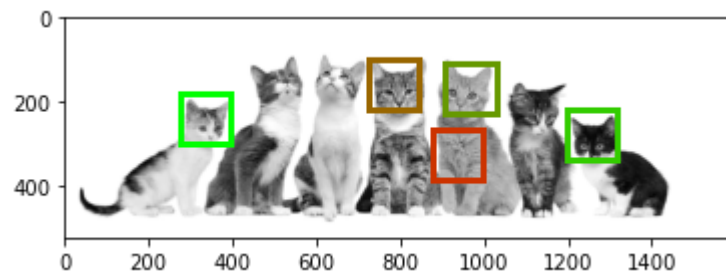


In [21]:

```
1 bsize=8
2 tsize=np.array([15,15]) #height and width in blocks
3 tsize_pix = bsize*tsize #height and width in pixels
4 posfiles = ('cat1.jpg','cat2.jpg','cat3.jpg','cat4.jpg','cat5.jpg')
5 negfiles = ('neg1.jpg','neg2.jpg','neg3.jpg','neg4.jpg','neg5.jpg')
6
7 # call learn_template to learn and visualize the template and training data
8 t = learn_template(posfiles,negfiles,tsize=tsize)
9
10 # call detect on one or more test images, visualizing the result with the plot_detections function
11 im = plt.imread("images/cats/cats1.jpg")
12 im = np.mean(im, axis=2)
13
14 detections = detect(im, t, 5)
15 plot_detections(im,detections,tsize_pix)
```









The detector works well when there is little interference and the object you are trying to detect has the same structure and angle as your

test images. It relies on similar magnitudes and orientation so images that have similar edge detection will create a more accurate template for detection. The test images would need to have a somewhat similar structure to the training images or else the detector would not categorize them well. Also, the detector didn't work well with images where edge detection was difficult. Sometimes images wouldn't have well defined edges and the functions would not consider them or there would be interference by the background of an image, which the template would consider for detection. This would create bias toward detection of images that were not intended.